

Lecture 7

Abstract Argumentation Frameworks

INST0074

Lecture Outline

- Abstract Argumentation Frameworks
 - Main Ideas & Definitions
- Acceptability Semantics
 - Extension-based semantics
 - Labelling-based semantics


AAFs: Main Ideas

- Dung, P.M. (1995). On the acceptability of arguments and its fundamental role in nonmonotonic reasoning, logic programming, and n-person games. *Artificial Intelligence*, 77:321-357, 1995.
- Arguments are defeasible entities that may attack each other
- The acceptance of an argument depends only on the status of the arguments that attack it.
- The structure, the origin and any other information about the arguments are abstracted away.
- Acceptability semantics formally define which arguments are accepted and which are rejected.

AAFs: Definitions

- An **argumentation framework** is a directed graph, the nodes of which are **arguments**, whereas the edges represent **attacks** among the arguments.
- $AF = \{A, R\}, R \subseteq A \times A$
 - A is a set of arguments
 - R is a binary relation on A
 - If $(a, b) \in R$ then we say that a **attacks** b
 - A set of arguments $S \subseteq A$ attacks an argument $b \in A$ iff there is an argument $a \in S$ that attacks b
 - A set of arguments $S \subseteq A$ is **conflict-free** iff there are no arguments $a, b \in S$ such that a attacks b

AAFs: An example



Argumentation Framework

$AF = \{A, R\}$

$A = \{a, b, c\}$

$R = \{(a,b), (b,a), (c,b)\}$

Conflict-free sets of arguments

$\{\}, \{a\}, \{b\}, \{c\}, \{a,c\}$

Attacks by sets of arguments

$\{a\}$ attacks b

$\{b\}$ attacks a

$\{c\}$ attacks b

$\{a, b\}$ attacks a

$\{a, b\}$ attacks b

$\{a, c\}$ attacks b

$\{b, c\}$ attacks a

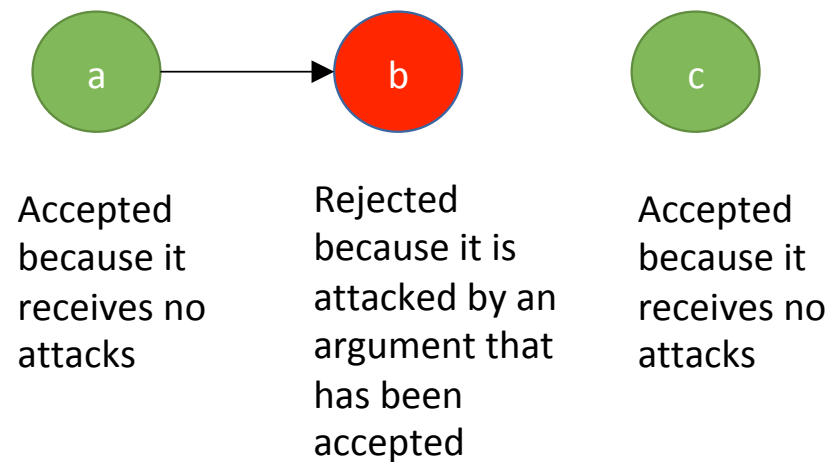
$\{b, c\}$ attacks b

$\{a, b, c\}$ attacks a


$\{a, b, c\}$ attacks b

Evaluation of arguments

- An argument is accepted if it does not receive any attacks.
- An argument is rejected if there is a counter-argument that has been accepted.
- An argument that does not attack and is not attacked by any other argument does not affect the acceptability of the other arguments.



A more complex case



- Arguments that are in conflict cannot be both accepted
- Should we accept neither or either of them?

- Scenario 1:

- a: The weather in Cuba is great, let's go there for our holidays.
- b: The tickets to Cuba are expensive, let's go somewhere else.

Accept either

- Scenario 2:

- a: Alice: Bob committed the murder. I was him in the crime scene.
- b: Bob: I didn't do it. Alice did it. She hated the victim!

Accept neither


Extension-based acceptability semantics

- The acceptability of arguments can be defined using the notion of extensions.
- An **extension** of an argumentation framework $AF = \{A, R\}$ is a set of arguments $E \subseteq A$ that we can reasonably accept.
- An **extension-based semantics** provides a formal way for identifying extensions (i.e. selecting sets of arguments that are reasonable to accept), according to some criterion.

Admissibility

- The notion of admissible sets of arguments can be regarded as the minimum requirement for a set of arguments to be accepted.
- A set of arguments $S \subseteq A$ **defends an argument** $a \in A$ iff it attacks any argument $b \in A$ that attacks a
- A set of arguments $E \subseteq A$ is **admissible** iff it is conflict-free and defends all its elements.

Admissibility (example)



Conflict-free	Admissible
{}	✓
{a}	✓
{b}	✗
{c}	✓
{a,c}	✓

It receives no attacks

It defends itself from b

It doesn't defend itself from c


It receives no attacks

It defends itself from b

Complete semantics

- Complete semantics is based on the notion of admissibility
 - A complete extension must be an admissible set of arguments
- It additionally requires accepting any argument that can be defended by an admissible set of arguments
- A set of arguments $E \subseteq A$ is a **complete extension** of $AF = \{A, R\}$ iff it is admissible and contains all the arguments it defends

Complete semantics (example)



Conflict-free	Admissible	Complete
{}	✓	✗
{a}	✓	✗
{b}	✗	✗
{c}	✓	✗
{a,c}	✓	✓

It defends c but doesn't contain it

It defends c but doesn't contain it

It is not admissible

It defends a but doesn't contain it

It contains all the arguments it defends


Grounded semantics

- The most conservative (sceptical) semantics regarding the number of arguments it accepts.
- It accepts only the arguments we cannot avoid to accept
- A set of arguments $E \subseteq A$ is a **grounded extension** of $AF = \{A, R\}$ iff it is the minimal (w.r.t. set inclusion) complete extension of AF

Minimal and maximal sets

- If S is set of sets
 - A set $X \in S$ is minimal iff there is no set $Y \in S$ such that $Y \subset X$
 - A set $X \in S$ is maximal iff there is no set $Y \in S$ such that $X \subset Y$
- For example, if $S = \{\{e\}, \{a,b\}, \{a,c\}, \{a,b,c\}, \{a,b,d\}, \{a,b,e\}, \{a,b,c,e\}\}$
 - The minimal sets are: $\{e\}, \{a,b\}, \{a,c\}$
 - The maximal sets are: $\{a,b,d\}, \{a,b,c,e\}$

Grounded semantics (example)




Conflict-free	Admissible	Complete	Grounded
{}	✓	✗	✗
{a}	✓	✗	✗
{b}	✗	✗	✗
{c}	✓	✗	✗
{a,c}	✓	✓	✓

The only complete extension
is also a grounded extension

Preferred semantics

- The most credulous semantics.
- It accepts as many arguments as possible
- A set of arguments $E \subseteq A$ is a **preferred extension** of $AF = \{A, R\}$ iff it is a maximal (w.r.t. set inclusion) complete extension of AF

Preferred semantics (example)




Conflict-free	Admissible	Complete	Grounded	Preferred
{}	✓	✗	✗	✗
{a}	✓	✗	✗	✗
{b}	✗	✗	✗	✗
{c}	✓	✗	✗	✗
{a,c}	✓	✓	✓	✓

The only complete extension
is also a preferred extension

Stable semantics

- It requires that every argument is either accepted or attacked by an accepted argument (and is therefore rejected).
- A set of arguments $E \subseteq A$ is a **stable extension** of $AF = \{A, R\}$ iff it is conflict-free and attacks all arguments in $A \setminus E$


Stable semantics (example)



Conflict-free	Admissible	Complete	Grounded	Preferred	Stable
{}	✓	✗	✗	✗	✗
{a}	✓	✗	✗	✗	✗
{b}	✗	✗	✗	✗	✗
{c}	✓	✗	✗	✗	✗
{a,c}	✓	✓	✓	✓	✓


It attacks the arguments it doesn't contain.

More examples




Admissible	Complete	Grounded	Preferred	Stable
{}	✓	✓	✗	✗
{a}	✓	✗	✓	✓
{b}	✓	✗	✓	✓

More examples



Admissible	Complete	Grounded	Preferred	Stable
{}	✗	✗	✗	✗
{a}	✓	✓	✗	✗
{a,c}	✓	✗	✓	✗
{a,d}	✓	✗	✓	✓
{c}	✗	✗	✗	✗
{d}	✗	✗	✗	✗

More examples



Admissible	Complete	Grounded	Preferred	Stable
{}	✓	✓	✗	✗
{d}	✗	✗	✗	✗
{a,d}	✓	✗	✓	✓

Properties of extensions

- The empty set is always admissible.
- There is always a preferred extension.
- The grounded extension is the intersection of all complete extensions and is unique.
- No stable extension is empty but there are argument frameworks for which there is no stable extension.
 - Consider for example this: $A = \{a, b, c\}$, $R = \{(a,b), (b,c), (c,a)\}$
- Every stable extension is also a preferred extension.
- If an argument graph has no cycle then there is a single extension that is stable, preferred, complete and grounded.

Labelling-based acceptability semantics

- Each argument in the framework is assigned a **label**:
 - $\text{Lab}(a) = \text{in}$: the argument is accepted
 - $\text{Lab}(a) = \text{out}$: the argument is rejected
 - $\text{Lab}(a) = \text{undec}$: the argument is neither accepted nor rejected
- A **labelling-based semantics** provides a way to select “reasonable” labellings among all the possible ones, according to some criterion.
- Legal labellings
 - An **in**-labelled argument is said to be **legally in** iff all its attackers are labelled **out**.
 - An **out**-labelled argument is said to be **legally out** iff at least one of its attackers is labelled **in**.
 - An **undec**-labelled argument is said to be **legally undec** iff not all its attackers are labelled **out** and it doesn't have an attacker that is labelled **in**.

Conflict-free and admissible labellings

- Let $AF = \{A, R\}$ be an argumentation framework and a labelling **Lab**
- **Lab** is an **admissible labelling** iff
 - every **in**-labelled argument is legally **in** and
 - every **out**-labelled argument is legally **out**.
- **Lab** is **conflict-free** iff for each argument $a \in A$ it holds that:
 - If a is labelled **in** then it does not have an attacker that is labelled **in**.
 - If a is labelled **out** then it has at least one attacker that is labelled **in**.
- Every admissible labelling is conflict-free.

Complete labellings

- **Lab** is a **complete labelling** iff
 - every **in**-labelled argument is legally **in**,
 - every **out**-labelled argument is legally **out** and
 - every **undec**-labelled argument is legally **undec**


An equivalent definition

- **Lab** is a **complete labelling** iff for each argument $a \in A$ it holds that:
 - a is labelled **in** iff all its attackers are labelled **out**.
 - a is labelled **out** iff it has at least one attacker that is labelled **in**.
- Every complete labelling is admissible.


Grounded, preferred and stable labellings

- **Lab** is a **grounded labelling** iff it is complete and **in(Lab)** (the set of arguments that are labelled **in**) is minimal among the complete labellings.
- **Lab** is a **preferred labelling** iff it is complete and **in(Lab)** is maximal among the complete labellings.
- **Lab** is a **stable labelling** iff it is complete and no argument is labelled **undec**.

An example



An example




Lab(a) = in

Lab(b) = in

Conflict-free	✗
Admissible	✗
Complete	✗
Grounded	✗
Preferred	✗
Stable	✗

An example




Lab(a) = out

Lab(b) = out

Conflict-free	✗
Admissible	✗
Complete	✗
Grounded	✗
Preferred	✗
Stable	✗

An example




Lab(a) = in

Lab(b) = out

Conflict-free	✓
Admissible	✓
Complete	✓
Grounded	✗
Preferred	✓
Stable	✓

An example




Lab(a) = out

Lab(b) = in

Conflict-free	✓
Admissible	✓
Complete	✓
Grounded	✗
Preferred	✓
Stable	✓

An example




Lab(a) = in

Lab(b) = undec

Conflict-free	✓
Admissible	✗
Complete	✗
Grounded	✗
Preferred	✗
Stable	✗

An example




Lab(a) = undec

Lab(b) = in

Conflict-free	✓
Admissible	✗
Complete	✗
Grounded	✗
Preferred	✗
Stable	✗

An example




Lab(a) = out

Lab(b) = undec

Conflict-free	✗
Admissible	✗
Complete	✗
Grounded	✗
Preferred	✗
Stable	✗

An example




Lab(a) = undec

Lab(b) = out

Conflict-free	✗
Admissible	✗
Complete	✗
Grounded	✗
Preferred	✗
Stable	✗

An example



Lab(a) = undec

Lab(b) = undec

Conflict-free	✓
Admissible	✓
Complete	✓
Grounded	✓
Preferred	✗
Stable	✗

Relation between extensions and labellings

- There is one-to-one relationship between most types of labellings and their corresponding extensions. Consider the following two functions:
- A function that takes as input a labelling **Lab** and returns the set of arguments that are **in**:

$$\text{Lab2Ext}(\text{Lab}) = \text{in}(\text{Lab})$$

- A function that take as input an extension **E** and returns a set of labelled arguments:

$$\text{Ext2Lab}(E) = \{(a, \text{in}) | a \in E\} \cup \{(b, \text{out}) | b \text{ is attacked by } E\} \cup \{(c, \text{undec}) | c \notin E \text{ and } c \text{ is not attacked by } E\}$$

- Then for the complete, preferred, grounded and stable semantics it holds:
 - If **E** is a **co/pr/gr/st** extension then **Ext2Lab(E)** is a **co/pr/gr/st** labelling
 - If **Lab** is a **co/pr/gr/st** labelling then **Lab2Ext(Lab)** is a **co/pr/gr/st** extension

Summing up

- Abstract Argumentation Frameworks is a **simple** but **powerful** model of arguments and argumentation-based inference.
 - **Simple:** It treats arguments as atomic entities (without an internal structure) and uses a single binary relation to model any type of attack.
 - **Powerful:** It enables many different ways of assessing the acceptability of arguments (acceptability semantics), each implementing a different form of non-monotonic reasoning.
 - It has been shown that several non-monotonic logics (Default Logic, Defeasible Logic, Logic Programming with negation as failure, etc.) are instances of Abstract Argumentation Frameworks.